

Table 1: Reproduction vs the paper’s Table 4 (F1). Each model occupies two rows: our reproduced score and, beneath it, $\Delta = \text{ours} - \text{paper}$. The Avg column averages over the 8/8 columns present on both sides.

Model	VLGuard	JailBreakV	SPA-VL	MMDS-Q	MMDS-R	VLSBench	MM-Safety	BeaverTails-V	Avg
Qwen3-VL-8B	0.6003	0.7593	0.6214	–	–	0.3207	0.5005	0.7273	0.5882
Δ vs paper	+0.0468	-0.1691	-0.0786	–	–	-0.4377	-0.1401	+0.0766	-0.0611
Qwen3-VL-4B	0.8238	0.8297	0.6824	–	–	0.8758	0.8820	0.5611	0.7758
Δ vs paper	+0.2536	-0.0909	+0.0319	–	–	+0.1890	+0.2717	-0.1617	+0.1389
LlavaGuard	0.4290	0.4756	–	–	–	–	–	–	0.4523
Δ vs paper	-0.0713	+0.1566	–	–	–	–	–	–	-0.0417
GuardReasoner-VL-7B	0.8822	0.9005	0.6541	–	–	0.6130	0.6318	0.5149	0.6994
Δ vs paper	-0.0211	-0.0447	-0.0443	–	–	-0.0176	-0.0380	-0.0271	+0.0510
SingGuard-2B	0.9549	0.9154	0.6960	0.9379	0.9291	0.9889	0.9530	0.8071	0.8978
Δ vs paper	+0.0158	-0.0351	-0.0822	-0.0318	+0.0331	+0.0286	+0.1121	+0.0024	+0.0054
SingGuard-4B	0.9472	0.9071	0.7005	0.9032	0.9046	0.9712	0.7891	0.8427	0.8707
Δ vs paper	+0.0176	-0.0623	-0.0755	-0.0573	+0.0317	+0.0179	-0.0714	+0.0090	-0.0238
SingGuard-8B	0.9536	0.9083	0.7089	0.9571	0.8908	0.9551	0.8653	0.8234	0.8828
Δ vs paper	+0.0025	-0.0645	-0.0795	-0.0280	-0.0055	-0.0117	-0.0159	-0.0081	-0.0263

Notes

- **Configuration.** Released checkpoints; SingGuard rows use `thinking_type=fast` (the paper’s Sec. 2.6 mode for high-throughput benchmark evaluation), greedy decoding, `max_model_len=32768`.
- **SPA-VL protocol.** The official test sets are EvalHarm (265) + EvalHelp (265); their filenames are *not* query-safety labels (EvalHelp measures helpfulness and contains hate-annotated images; EvalHarm measures refusal behaviour). We map `harm`→`unsafe` / `help`→`safe`, the reading most consistent with the paper’s numbers for this column; treat its absolute value as protocol-dependent.
- **JailBreakV coverage.** HF hosts only ~360 of the 28,000 images, so 662/1000 of our sampled rows run text-only. On the image-bearing subset (the benchmark’s designed form) SingGuard-8B scores F1 0.9600 vs the paper’s 0.9728.
- **MMDS protocol.** The corpus carries an official `set` field (train 4045 / val 109 / test 330); we score the *test* split, as the paper does. Scoring a stratified sample of the whole corpus instead drops SingGuard-8B to 0.7722/0.7209 — the earlier gap in this column was a sampling error on our side, not a protocol or checkpoint difference.
- **MM-Safety is all-unsafe.** The benchmark is a set of malicious instructions (every sampled row is gold-unsafe), so its F1 reduces to a precision-weighted recall on the attack set and is sensitive to which attacks are sampled.
- **Baseline harness fidelity.** Each baseline is driven through its own released prompt/template, not a common one. Two defects found and fixed this way: LlavaGuard’s template ships as `chat_template.json` with no inline copy, so vLLM silently substituted a generic template and the model returned JSON missing its leading keys on 26% of samples; and GuardReasoner-VL’s official interface (INSTRUCTION as a system message, a `Human user:/AI assistant:` transcript, verdict in a `<result>` block) differs from a plain chat call. Baseline prompts are still approximations in general — the paper publishes only SingGuard’s templates — so their per-column deltas are indicative, not exact.
- **Still missing.** LlamaGuard3-Vision-11B and LlamaGuard4-12B rows (queued); LLaVASHield-v1.0-7B ships a LLaVA-NeXT training-format checkpoint (LlavaQwenForCausalLM, no `auto_map`) that neither vLLM 0.11 nor transformers loads directly — a converted copy is being validated.

- **Qwen3-VL-235B** is not run: bf16 (~ 470 GB) / FP8 (~ 235 GB) does not fit a shared $8\times A6000$ node without multi-GPU tensor parallelism.
- Baseline prompts are approximations (the paper publishes only SingGuard’s templates); per-column deltas for baseline rows are therefore indicative, not exact.