

Table 1: Reproduction vs the paper’s Table 4 (F1). Each model occupies two rows: our reproduced score and, beneath it, $\Delta = \text{ours} - \text{paper}$. The Avg column averages over the 6/8 columns present on both sides.

oprule Model	VLGuard	JailBreakV	SPA-VL	VLSBench	MM-Safety	BeaverTails-V	Avg
Qwen3-VL-8B	0.6003	0.7593	0.6214	0.3207	0.5005	0.7273	0.5883
Δ vs paper	+0.0468	-0.1691	-0.0786	-0.4377	-0.1401	+0.0766	-0.1170
Qwen3-VL-4B	0.8238	0.8297	0.6824	0.8758	0.8820	0.5611	0.7758
Δ vs paper	+0.2536	-0.0909	+0.0319	+0.1890	+0.2717	-0.1617	+0.0823
SingGuard-8B	0.9536	0.9083	0.7089	0.9551	0.8653	0.8234	0.8691
Δ vs paper	+0.0025	-0.0645	-0.0795	-0.0117	-0.0159	-0.0081	-0.0295