

Table 1: Reproduction of the SingGuard technical report’s Table 4: multimodal safety benchmark results (F1 on the unsafe class) using the released checkpoints in `thinking_type=fast` with greedy decoding. MMDS columns are still running; Columns present: 6/8.

Model	VLGuard	JailBreakV	SPA-VL	MMDS-Q	MMDS-R	VLSBench	MM-Safety	BeaverTails-V	Avg
Qwen3-VL-8B	0.6003	0.7593	0.6214	–	–	0.3207	0.5005	0.7273	0.5882
Qwen3-VL-4B	0.8238	0.8297	0.6824	–	–	0.8758	0.8820	0.5611	0.7758
SingGuard-4B	0.9472	0.9071	0.7005	–	–	0.9712	0.7891	0.8427	0.8596
SingGuard-8B	0.9536	0.9083	0.7089	–	–	0.9551	0.8653	0.8234	0.8691

Notes

- **Configuration.** Released checkpoints; SingGuard rows use `thinking_type=fast` (the paper’s Sec. 2.6 mode for high-throughput benchmark evaluation), greedy decoding, `max_model_len=32768`.
- **SPA-VL protocol.** The official test sets are EvalHarm (265) + EvalHelp (265); their filenames are *not* query-safety labels (EvalHelp measures helpfulness and contains hate-annotated images; EvalHarm measures refusal behaviour). We map harm→unsafe / help→safe, the reading most consistent with the paper’s numbers for this column; treat its absolute value as protocol-dependent.
- **JailBreakV coverage.** HF hosts only ~360 of the 28,000 images, so 662/1000 of our sampled rows run text-only. On the image-bearing subset (the benchmark’s designed form) SingGuard-8B scores F1 0.9600 vs the paper’s 0.9728.
- **Pending.** MMDS-Q/MMDS-R columns (data ready, runs in flight); baseline guards (LLaVASHield-v1.0-7B, LlamaGuard3-Vision-11B, LlamaGuard4-12B, ShieldGemma-2-4B, SafeGuard-VL-RL, GuardReasoner-VL-7B, LlavaGuard-v1.0-7B) are downloading.
- **Qwen3-VL-235B** is not run: bf16 (~470 GB) / FP8 (~235 GB) does not fit a shared 8×A6000 node without multi-GPU tensor parallelism.
- Baseline prompts are approximations (the paper publishes only SingGuard’s templates); per-column deltas for baseline rows are therefore indicative, not exact.